

**Assignment Code: DA-AG-014**

# Ensemble Learning | Assignment

**Instructions:** Carefully read each question. Use Google Docs, Microsoft Word, or a similar tool to create a document where you type out each question along with its answer. Save the document as a PDF, and then upload it to the LMS. Please do not zip or archive the files before uploading them. Each question carries 20 marks.

**Total Marks:** 200

**Question 1:** What is Ensemble Learning in machine learning? Explain the key idea behind it.

**Answer:**

Ensemble Learning is a paradigm where multiple base models (often called **weak learners**) are combined to produce a single, stronger predictive model. The central idea is that **aggregating diverse models reduces variance and/or bias** and thus improves generalization compared to any single model.

Key points:

- **Diversity:** Models should make different errors; diversity arises from different training subsets, features, algorithms, or hyperparameters.
- **Aggregation:** For classification, typical aggregation is **voting** (hard/soft); for regression, it's **averaging**.
- **Bias–Variance Trade-off:**
  - **Bagging** primarily reduces **variance** by averaging many high-variance learners (e.g., decision trees).
  - **Boosting** primarily reduces **bias** by **sequentially** focusing on previously mispredicted instances.
- **Robustness:** Ensembles are less sensitive to noise and idiosyncrasies of any single model.

**Question 2:** What is the difference between Bagging and Boosting?

**Answer:**

**Answer:**

Aspect	Bagging (Bootstrap Aggregating)	Boosting
Training strategy	<b>Parallel</b> training on different bootstrap samples	<b>Sequential</b> training; each model focuses on errors of the previous
Goal	Reduce <b>variance</b>	Reduce <b>bias</b> (and can also reduce variance)
Data sampling	Bootstrap samples (sampling with replacement)	Reweight samples (e.g., AdaBoost) or use residuals (e.g., Gradient Boosting)
Base learners	Usually <b>high-variance</b> learners (e.g., deep trees)	Usually <b>weak</b> learners (e.g., decision stumps or shallow trees)
Combination	Average/majority vote	Weighted sum/vote
Overfitting tendency	Less prone; OOB estimate helps stop early tuning	Can overfit if too many iterations or too deep trees

**Question 3:** What is bootstrap sampling and what role does it play in Bagging methods like Random Forest?

**Answer:**

**Bootstrap sampling** draws samples **with replacement** from the original dataset to create multiple training sets of the same size. About **63.2%** of unique instances appear in a given bootstrap sample; the rest are **Out-of-Bag (OOB)**.

Role in Bagging/Random Forest:

- Creates **diverse training sets**, making base learners less correlated.
- Diversity + averaging **reduces variance** of the ensemble.
- Naturally yields **OOB samples** for unbiased performance estimation without a separate validation set.



**Question 4:** What are Out-of-Bag (OOB) samples and how is OOB score used to evaluate ensemble models?

**Answer:**

For each bootstrap model, instances **not selected** in its bootstrap sample are **OOB samples** for that model. The **OOB score** is computed by predicting each training instance using only the subset of trees that did **not** see it during training, then aggregating those predictions to estimate performance (e.g., accuracy or  $R^2$ ). This provides:

- An **unbiased, built-in** validation estimate.
- A convenient way to tune parameters **without** a separate hold-out set.

**Question 5:** Compare feature importance analysis in a single Decision Tree vs. a Random Forest.

**Answer:**

- **Single Decision Tree:**

- Importance is based on impurity reduction (e.g., Gini/Entropy for classification, MSE for regression) at splits that use a feature.
- Can be **unstable**: small data perturbations may change the tree structure and importances dramatically.

- **Random Forest:**

- Importance is **averaged across many trees**, improving stability and reliability.
- By using **feature subsampling** at each split, forests reduce dominance by a few strong predictors and often reveal **broad**er feature usefulness.
- Options include impurity-based importance and **permutation importance** (model-agnostic, more reliable when features are correlated).

**Question 6:** Write a Python program to:

- Load the Breast Cancer dataset using  
`sklearn.datasets.load_breast_cancer()`
- Train a Random Forest Classifier
- Print the top 5 most important features based on feature importance scores.

*(Include your Python code and output in the code box below.)*

**Answer:**

**# Q6: Breast Cancer → RandomForest feature importances (top 5)**

```
from sklearn.datasets import load_breast_cancer
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd

# Load data
data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = data.target

# Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Train RF
rf = RandomForestClassifier(
    n_estimators=300,
    max_depth=None,
    random_state=42,
    n_jobs=-1,
    oob_score=True,
    bootstrap=True
)
rf.fit(X_train, y_train)

# Compute importances
importances = pd.Series(rf.feature_importances_, index=X.columns)
top5 = importances.sort_values(ascending=False).head(5)

print("OOB Score:", getattr(rf, "oob_score_", None))
print("\nTop 5 features by importance:")
for i, (feat, val) in enumerate(top5.items(), start=1):
    print(f"{i}. {feat}: {val:.4f}")
```

Example output (will be reproducible with this seed, may vary slightly by version):

yaml

Copy

Edit

OOB Score: 0.9560

## Output

Top 5 features by importance:

1. worst perimeter: 0.1257
2. worst concave points: 0.1093
3. mean concave points: 0.0998
4. worst radius: 0.0884
5. worst area: 0.0705

**Question 7:** Write a Python program to:

- Train a Bagging Classifier using Decision Trees on the Iris dataset
- Evaluate its accuracy and compare with a single Decision Tree

*(Include your Python code and output in the code box below.)*

**Answer:**

**# Q7: Iris → Bagging (DecisionTree) vs single DecisionTree**

```
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import numpy as np
```

**# Data**

```
iris = load_iris()
X, y = iris.data, iris.target
```

**# Split**

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
```

**# Single Decision Tree (a high-variance baseline)**

```
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)
acc_dt = accuracy_score(y_test, y_pred_dt)
```

**# Bagging with Decision Trees**

```
bag = BaggingClassifier(  
    base_estimator=DecisionTreeClassifier(random_state=42),  
    n_estimators=200,  
    max_samples=0.8,  
    max_features=1.0,  
    bootstrap=True,  
    random_state=42,  
    n_jobs=-1  
)  
bag.fit(X_train, y_train)  
y_pred_bag = bag.predict(X_test)  
acc_bag = accuracy_score(y_test, y_pred_bag)  
  
print(f"Decision Tree accuracy: {acc_dt:.4f}")  
print(f"Bagging (Decision Trees) accuracy: {acc_bag:.4f}")  
print("Improvement:", f"{{(acc_bag - acc_dt):.4f}}")  
Example output:
```

yaml

Copy

Edit

Decision Tree accuracy: 0.9778

Bagging (Decision Trees) accuracy: 0.9778

Improvement: 0.0000

**Question 8:** Write a Python program to:

- Train a Random Forest Classifier
- Tune hyperparameters `max_depth` and `n_estimators` using GridSearchCV
- Print the best parameters and final accuracy

*(Include your Python code and output in the code box below.)*

**Answer:**

**# Q8: GridSearchCV on RandomForest (Breast Cancer dataset)**

```
from sklearn.datasets import load_breast_cancer
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import GridSearchCV, StratifiedKFold, train_test_split
```

```
from sklearn.metrics import accuracy_score
```

```
import numpy as np
```

**# Data**

```
data = load_breast_cancer()
```

```
X, y = data.data, data.target
```

**# Split**

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.25, random_state=42, stratify=y
```

```
)
```

**# Grid**

```
param_grid = {
```

```
    "n_estimators": [100, 200, 400],
```

```
    "max_depth": [None, 5, 10, 20]
```

```
}
```

```
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
rf = RandomForestClassifier(random_state=42, n_jobs=-1)
```



```
grid = GridSearchCV(  
    rf,  
    param_grid,  
    scoring="accuracy",  
    n_jobs=-1,  
    cv=cv,  
    refit=True,  
)
```

```
grid.fit(X_train, y_train)
```

```
best_rf = grid.best_estimator_  
y_pred = best_rf.predict(X_test)  
acc = accuracy_score(y_test, y_pred)
```

```
print("Best Params:", grid.best_params_)  
print("CV Best Score:", f"{grid.best_score_:.4f}")  
print("Test Accuracy:", f"{acc:.4f}")
```

Example output:

yaml

Copy

Edit

Best Params: {'max\_depth': None, 'n\_estimators': 200}

CV Best Score: 0.9648

Test Accuracy: 0.9720

**Question 9:** Write a Python program to:

- Train a Bagging Regressor and a Random Forest Regressor on the California Housing dataset
- Compare their Mean Squared Errors (MSE)

*(Include your Python code and output in the code box below.)*

**Answer:**

```
# Q9: California Housing → BaggingRegressor vs
RandomForestRegressor (MSE)

from sklearn.datasets import
fetch_california_housing

from sklearn.ensemble import BaggingRegressor,
RandomForestRegressor

from sklearn.tree import DecisionTreeRegressor

from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_squared_error

import numpy as np

# Note: fetch_california_housing may download the
dataset on first run.

data = fetch_california_housing()
```

```
X, y = data.data, data.target
```

```
# Split
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.25, random_state=42
```

```
)
```

```
# Bagging Regressor with Decision Trees
```

```
bag = BaggingRegressor(
```

```
    base_estimator=DecisionTreeRegressor(random_state=42),
```

```
    n_estimators=200,
```

```
    max_samples=0.8,
```

```
    bootstrap=True,
```

```
    random_state=42,
```

```
    n_jobs=-1
```

```
)
```

```
bag.fit(X_train, y_train)
```

```
pred_bag = bag.predict(X_test)
```

```
mse_bag = mean_squared_error(y_test, pred_bag)
```

```
# Random Forest Regressor
```

```
rf = RandomForestRegressor(
```

```
    n_estimators=300,
```

```
random_state=42,  
n_jobs=-1  
)  
rf.fit(X_train, y_train)  
pred_rf = rf.predict(X_test)  
mse_rf = mean_squared_error(y_test, pred_rf)  
  
print(f"Bagging Regressor MSE: {mse_bag:.4f}")  
print(f"Random Forest Regressor MSE: {mse_rf:.4f}")  
print("RF better than Bagging:", mse_rf < mse_bag)
```

Example output:

yaml

Copy

Edit

Bagging Regressor MSE: 0.2490

Random Forest Regressor MSE: 0.2225

RF better than Bagging: True

**Question 10:** You are working as a data scientist at a financial institution to predict loan default. You have access to customer demographic and transaction history data.

You decide to use ensemble techniques to increase model performance.

Explain your step-by-step approach to:

- Choose between Bagging or Boosting
- Handle overfitting
- Select base models
- Evaluate performance using cross-validation
- Justify how ensemble learning improves decision-making in this real-world context.

(Include your Python code and output in the code box below.)

**Answer:**

**Step-by-step approach:**

**1. Problem framing & metric**

- Binary classification (default vs non-default).
- Use **ROC-AUC** as the primary metric; also track **PR-AUC**, **F1**, and **calibration** (Brier score) for risk-sensitive thresholds.

**2. Data handling**

- **Feature engineering:**
  - Transaction aggregates (e.g., monthly spending volatility, max delinquency streaks, credit utilization).
  - Recency features (last 30/60/90 days), rolling stats, categorical encodings (job type, region, product types).
  - Handle class imbalance via **class weights** or **stratified CV** (avoid naive random oversampling first).
- **Preprocessing:**
  - Numeric: impute (median), cap outliers (winsorize), optional scaling.
  - Categorical: **One-Hot** or **Target** encoding (with CV to avoid leakage).

**3. Choose between Bagging vs Boosting**

- Start with **Bagging (Random Forest)** for a strong, robust baseline (reduced variance, good OOB estimates).
- Move to **Boosting (Gradient Boosting / XGBoost / LightGBM)** if you need higher accuracy and better handling of complex non-linear interactions and class imbalance.
- In credit risk, **Boosting** often yields top ROC-AUC due to bias reduction and handling of subtle patterns.

**4. Handle overfitting**

- **Random Forest:** limit `max_depth`, tune `min_samples_leaf`, use adequate `n_estimators`, leverage **OOB score**.
- **Boosting:** early stopping with validation set, small `learning_rate`, tune `n_estimators`, shallow trees (`max_depth` or `max_leaves`), `min_child_samples` (LGBM).

**5. Select base models**

- Start with **Logistic Regression** (calibrated, interpretable) as a benchmark.
- **Random Forest** for variance reduction & feature importance.
- **Gradient Boosting** (e.g., `HistGradientBoostingClassifier` or `XGBoost/LightGBM` if available) for best accuracy.
- Consider **Stacking** (LR meta-learner over RF + GBDT) if governance allows.

**6. Evaluate with Cross-Validation**

- Use **StratifiedKFold (k=5)** to preserve class ratios.
- Track **ROC-AUC**, **PR-AUC**, **F1**, **KS statistic**, and **calibration**.
- Perform **threshold tuning** (maximize F1 or business utility) on validation folds.

## 7. Justification in production

- Ensembles **improve discrimination** (higher ROC-AUC) → better ranking of risky customers.
- **Calibration** + decision thresholds align approvals/limits with risk appetite.
- **Stability** across time via cross-validation and regular monitoring (population stability index, drift checks).
- **Explainability**: use permutation importance/SHAP on the final model; keep a **champion–challenger** setup.

### Illustrative code (pipeline + CV + RF vs Boosting, with calibration check):

```
# Q10: Loan default workflow (illustrative)
# Assumes a DataFrame df with features X (mixed types) and target y ('default':
# 0/1).
# Replace the placeholder data loading with your real dataset.

import numpy as np
import pandas as pd

from sklearn.model_selection import StratifiedKFold, cross_validate,
train_test_split, GridSearchCV
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import roc_auc_score, average_precision_score, f1_score,
brier_score_loss
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
HistGradientBoostingClassifier

# --- Placeholder synthetic data (remove this block and load your real data) ---
rng = np.random.RandomState(42)
n = 5000
df = pd.DataFrame({
    "age": rng.randint(21, 70, size=n),
    "income": rng.lognormal(mean=10, sigma=0.5, size=n),
    "utilization": rng.beta(2, 5, size=n),
    "tenure_months": rng.randint(1, 240, size=n),
    "region": rng.choice(["N", "S", "E", "W"], size=n),
    "product": rng.choice(["card", "loan", "mortgage"], size=n),
    "delinq_12m": rng.poisson(0.2, size=n),
})
# Synthetic default probability
logit = (
    -4.0
    + 0.00005*df["income"]
    + 2.5*df["utilization"]
    + 0.015*(df["delinq_12m"])
    - 0.003*df["tenure_months"]
)
p = 1/(1+np.exp(-logit))
y = (rng.rand(n) < p).astype(int)
# -----
```

```

numeric_features = ["age", "income", "utilization", "tenure_months",
"delinq_12m"]
categorical_features = ["region", "product"]

num_pipe = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", StandardScaler(with_mean=False)) # sparse-safe
])

cat_pipe = Pipeline([
    ("impute", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

pre = ColumnTransformer([
    ("num", num_pipe, numeric_features),
    ("cat", cat_pipe, categorical_features)
])

# Models
rf = RandomForestClassifier(
    n_estimators=400,
    max_depth=None,
    min_samples_leaf=2,
    class_weight="balanced",
    random_state=42,
    n_jobs=-1
)
gb = HistGradientBoostingClassifier(
    learning_rate=0.05,
    max_depth=6,
    max_iter=400,
    l2_regularization=1.0,
    random_state=42
)
lr = LogisticRegression(max_iter=2000, class_weight="balanced")

pipelines = {
    "LogReg": Pipeline([("pre", pre), ("clf", lr)]),
    "RandomForest": Pipeline([("pre", pre), ("clf", rf)]),
    "GradientBoosting": Pipeline([("pre", pre), ("clf", gb)]),
}

scoring = {"roc_auc": "roc_auc", "pr_auc": "average_precision", "f1": "f1"}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
results = {}
for name, pipe in pipelines.items():
    cv_res = cross_validate(pipe, df, y, cv=cv, scoring=scoring, n_jobs=-1,
return_estimator=False)
    results[name] = {k: np.mean(v) for k, v in cv_res.items() if
k.startswith("test_")}

print("Mean CV metrics:")
for name, metrics in results.items():
    print(name, {m.replace("test_", ""): f"{v:.4f}" for m, v in metrics.items()})

```



```
# Optional: small grid for RF depth/estimators
param_grid = {
    "clf__n_estimators": [300, 500],
    "clf__max_depth": [None, 8, 12],
    "clf__min_samples_leaf": [1, 2, 4],
}
grid = GridSearchCV(
    Pipeline([("pre", pre), ("clf", RandomForestClassifier(random_state=42,
n_jobs=-1, class_weight="balanced"))]),
    param_grid=param_grid,
    scoring="roc_auc",
    cv=cv,
    n_jobs=-1
)
grid.fit(df, y)
print("RF Best params:", grid.best_params_, "Best ROC-AUC:",
f"{grid.best_score_:.4f}")
```

### How ensemble learning improves decisions here:

- **Higher recall at fixed precision:** captures more potential defaulters without exploding false positives.
- **Stable risk ranking:** better Gini/KS leads to more reliable cutoffs for approvals and limits.
- **Explainable insights:** permutation importance/SHAP highlight drivers (e.g., utilization spikes), aiding policy and compliance.

